FAIR CLASSIFICATION

Explosion of fairness research over last five years

Fair classification is the most common setup, involving:

- X, some data
- Y, a label to predict
- $\hat{Y}$, the model prediction
- A, a sensitive attribute (race, gender, age, socio-economic status)

We want to learn a classifier that is:

- accurate
- fair with respect to A

REPRESENTATIONS BEYOND CLUSTERS

Aim: Replace discrete representation with continuous, multi-dimensional Z

Allow more flexible, nuanced representations

Bring ML arsenal to bear: powerful methods for mapping, embedding in vector spaces: Variational Auto Encoders (VAE)

How to maintain statistical parity in learned representations?

VAE

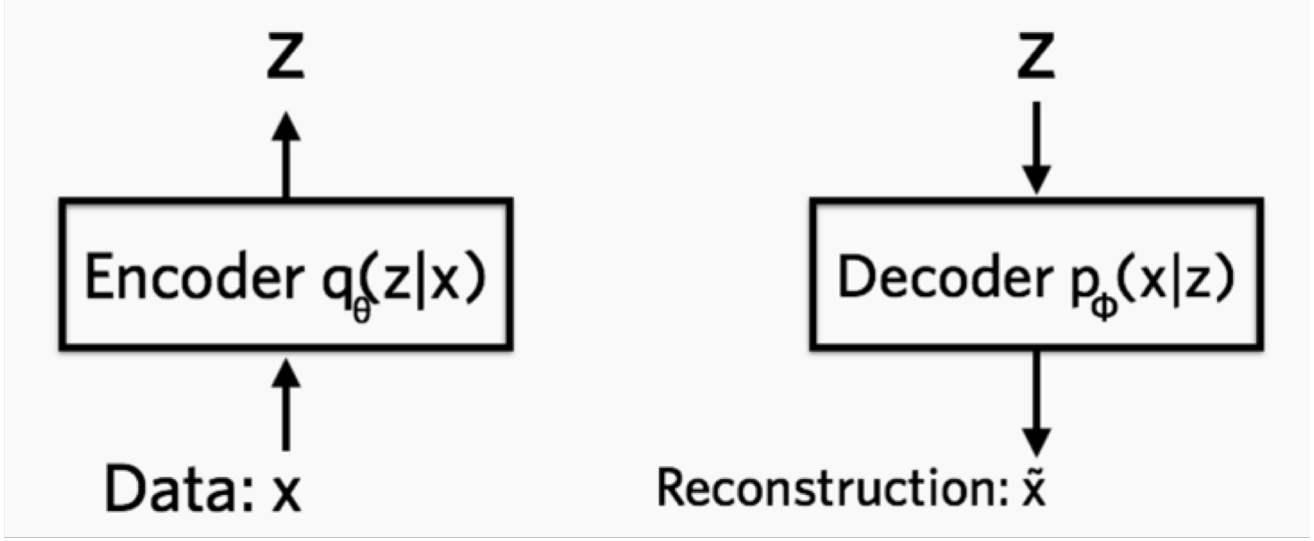

Re-formulation of autoencoders:

- Each input encoded into a distribution in latent space
- Output prediction obtained by sampling from distribution, mapping through decoder

Allows maximum-likelihood based density modelling:

$$\log p_\theta(\mathbf{x}) \geq \mathbb{E}_{q_\phi(\mathbf{z}|\mathbf{x})}\left[\log p_\theta(\mathbf{x}\,|\,\mathbf{z})\right] - D_{KL}\left(q_\phi(\mathbf{z}\,|\,\mathbf{x}) \,\|\, p(\mathbf{z})\right)$$

MMD

- Suppose we have access to samples from two probability distributions $X \sim P_A$ and $Y \sim P_B$, how can we tell if $P_A = P_B$?
- Maximum Mean Discrepancy (MMD) is a measure of distance between two distributions given only samples from each. [Gretton 2010]

$$\left\| \frac{1}{N} \sum_{n=1}^{N} \phi(X_n) - \frac{1}{M} \sum_{m=1}^{M} \phi(Y_m) \right\|^2$$

$$= \frac{1}{N^2} \sum_{n=1}^{N} \sum_{n'=1}^{N} \phi(X_n)^\top \phi(X_{n'}) + \frac{1}{M^2} \sum_{m=1}^{M} \sum_{m'=1}^{M} \phi(Y_m)^\top \phi(Y_{m'}) - \frac{2}{NM} \sum_{n=1}^{N} \sum_{m=1}^{M} \phi(X_n)^\top \phi(Y_m)$$

$$= \frac{1}{N^2} \sum_{n=1}^{N} \sum_{n'=1}^{N} k(X_n, X_{n'}) + \frac{1}{M^2} \sum_{m=1}^{M} \sum_{m'=1}^{M} k(Y_m, Y_{m'}) - \frac{2}{MN} \sum_{n=1}^{N} \sum_{m=1}^{N} k(X_n, Y_m)$$

- Our idea: learn to make two distributions indistinguishable

  ➔ small MMD!

VARIATIONAL FAIR AUTOENCODER

VAE with regularizer on latent representations

Match higher-order moments, continuous *Z*:

$$\ell_{\text{MMD}}(\mathbf{Z}_{1\mathbf{s}=0}, \mathbf{Z}_{1\mathbf{s}=1}) = \| \mathbb{E}_{\tilde{p}(\mathbf{x}|\mathbf{s}=0)}[\mathbb{E}_{q(\mathbf{z}_1|\mathbf{x},\mathbf{s}=0)}[\psi(\mathbf{z}_1)]] - E_{\tilde{p}(\mathbf{x}|\mathbf{s}=1)}[\mathbb{E}_{q(\mathbf{z}_1|\mathbf{x},\mathbf{s}=1)}[\psi(\mathbf{z}_1)]]\|^2$$

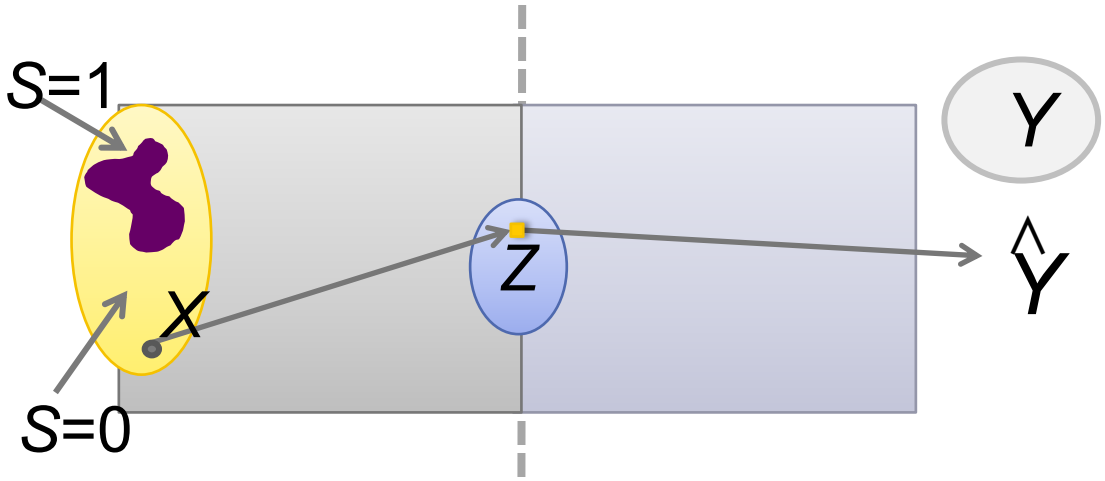

VARIATIONAL FAIR AUTOENCODER

Extend VAE to include some labels *y* (semi-supervised VAE [Kingma & Welling, 2014]) and "nuisance variable" *s*

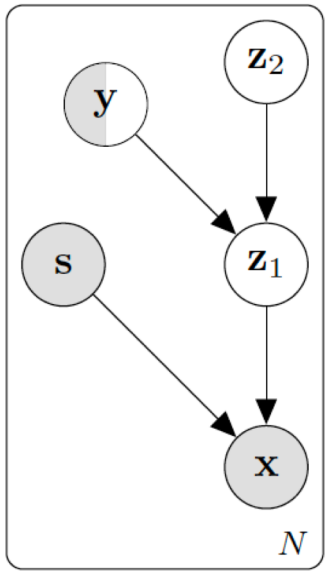

Objective -- maximize:

$$\sum_{n=1}^{N_s} \mathbb{E}_{q_\phi(\mathbf{z}_{1n}|\mathbf{x}_n,\mathbf{s}_n)}[-KL(q_\phi(\mathbf{z}_{2n}|\mathbf{z}_{1n},\mathbf{y}_n)||p(\mathbf{z}_2)) + \log p_\theta(\mathbf{x}_n|\mathbf{z}_{1n},\mathbf{s}_n)] +$$

$$+ \mathbb{E}_{q_\phi(\mathbf{z}_{2n}|\mathbf{z}_{1n},\mathbf{y}_n)}[-KL(q_\phi(\mathbf{z}_{1n}|\mathbf{x}_n,\mathbf{s}_n)||p_\theta(\mathbf{z}_{1n}|\mathbf{z}_{2n},\mathbf{y}_n))]$$

Add for labeled set:

$$\sum_{n=1}^{N} \mathbb{E}_{q(\mathbf{z}_{1n}|\mathbf{x}_n,\mathbf{s}_n)}[-\log q_\phi(\mathbf{y}_n|\mathbf{z}_{1n})]$$

unlabeled set:

$$\sum_{m=1}^{M} \mathbb{E}_{q_\phi(\mathbf{z}_{1m}|\mathbf{x}_m,\mathbf{s}_m)}[-KL(q(\mathbf{y}_m|\mathbf{z}_{1m})||p(\mathbf{y}_m))$$

RESULTS

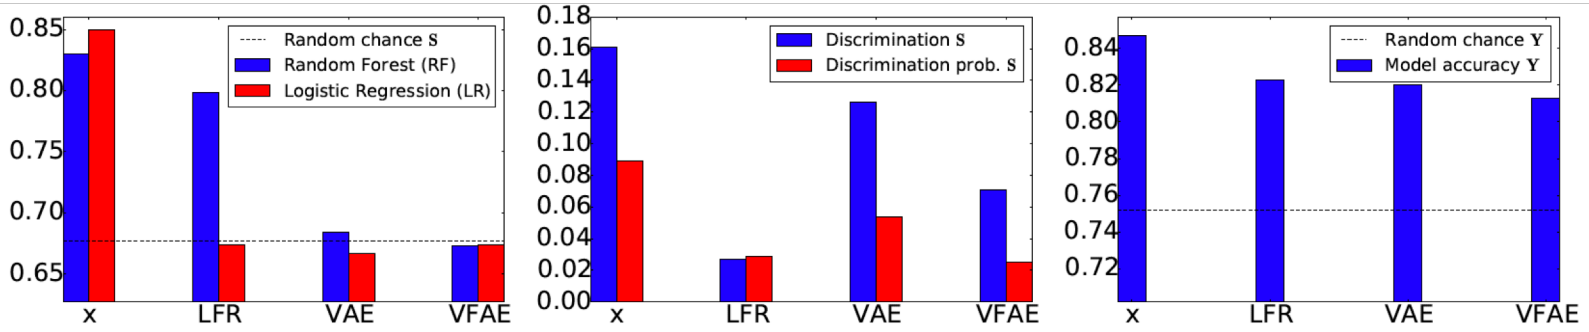

(a) Adult dataset

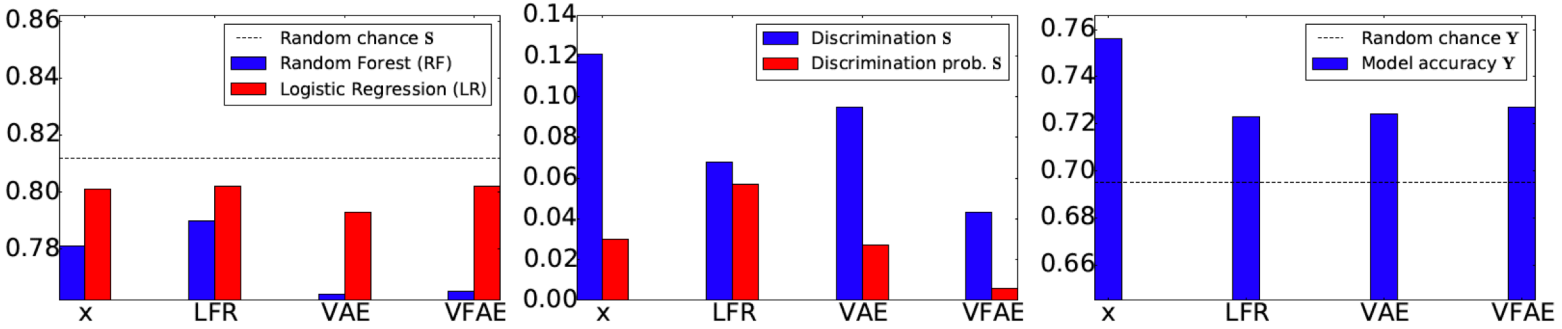

(b) German dataset

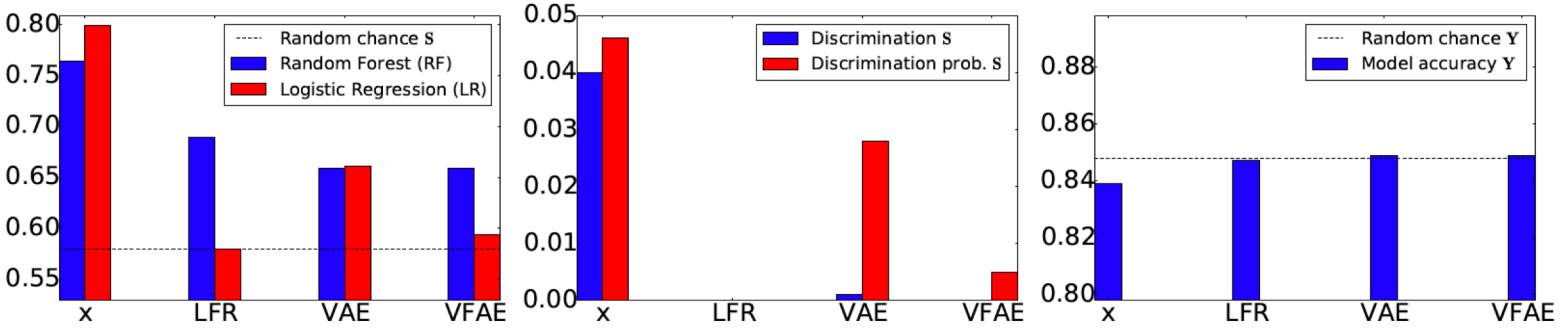

(c) Health dataset

RESULTS

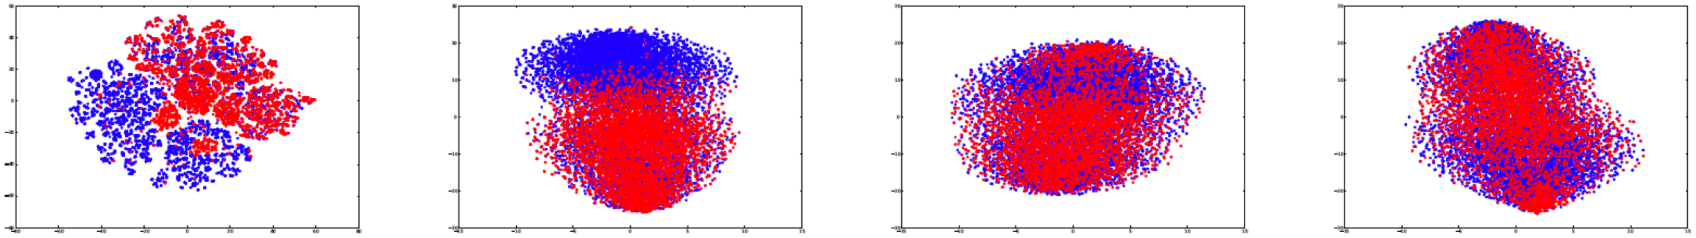

(a) (b) (c) (d)

Figure 4: t-SNE (van der Maaten, 2013) visualizations from the Adult dataset on: (a): original $\mathbf{x}$, (b): latent $\mathbf{z}_1$ without $\mathbf{s}$ and MMD, (c): latent $\mathbf{z}_1$ with $\mathbf{s}$ and without MMD, (d): latent $\mathbf{z}_1$ with $\mathbf{s}$ and MMD. Blue colour corresponds to males whereas red colour corresponds to females.

ADAPTING THE FRAMEWORK

The same idea has many other useful applications, e.g.,

- Eliminating demographic discrimination in deciding who should get transplant surgery
- Removing confounds, such as which scanner produced a medical image

Key: Learning to make two (or more) distributions indistinguishable

DOMAIN ADAPTATION

Natural fit: domain adaptation

Make feature representations for source and target domain data indistinguishable

Sentiment classification

- Product reviews (text, tf-idf on words & bigrams)
- Labeled data from source domain, unlabeled data from target domain

Source - Target	S		Y	
	RF	LR	VFAE	DANN
books - dvd	0.535	0.564	**0.799**	0.784
books - electronics	0.541	0.562	**0.792**	0.733
books - kitchen	0.537	0.583	**0.816**	0.779
dvd - books	0.537	0.563	**0.755**	0.723
dvd - electronics	0.538	0.566	**0.786**	0.754
dvd - kitchen	0.543	0.589	**0.822**	0.783
electronics - books	0.562	0.590	**0.727**	0.713
electronics - dvd	0.556	0.586	**0.765**	0.738
electronics - kitchen	0.536	0.570	0.850	**0.854**
kitchen - books	0.560	0.593	**0.720**	0.709
kitchen - dvd	0.561	0.599	0.733	**0.740**
kitchen - electronics	0.533	0.565	0.838	**0.843**

LEARNING INVARIANT FEATURES

If we have labeled data from all domains, factoring out unwanted domain bias still leads to better generalization.

Make the learned representations invariant to unwanted transformation / variation / bias.

Example: Face identification under different lighting conditions

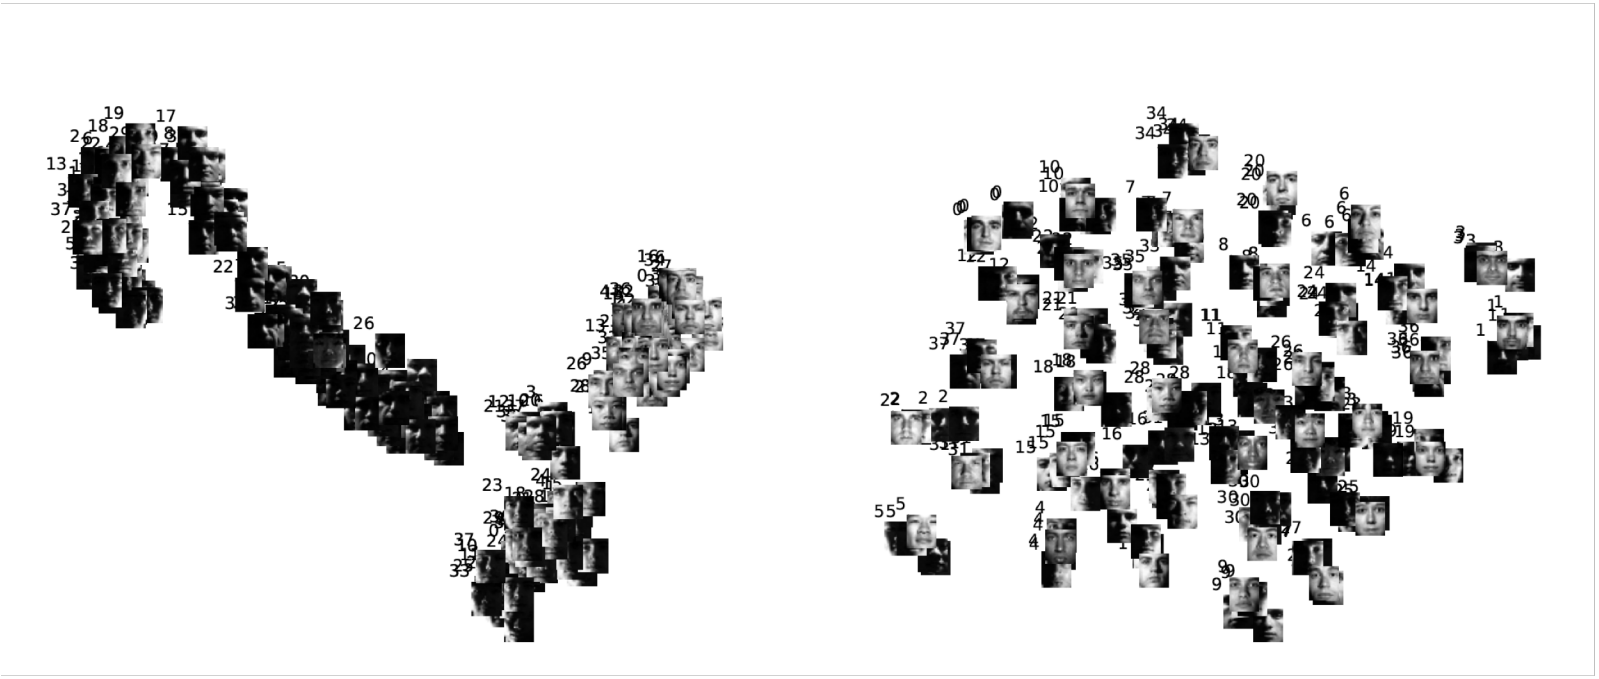

ADVERSARIAL FAIR LEARNING

Rather than using MMD to ensure learned representation is fair, can use adversarial approach

Adversary takes latent representation (here *R*) as input and attempts to predict *S*, then model minimizes:

$$D_{\theta,\phi}(R,S) = \mathop{\mathbb{E}}_{X,S} S \cdot \log\left(\mathrm{Adv}(R)\right) + (1-S) \cdot \log\left(1-\mathrm{Adv}(R)\right)$$

Combine with reconstruction and classification losses to ensure representation retains info about *X, Y*

$$C_\theta(X,R) = \mathop{\mathbb{E}}_{X} \|X - \mathrm{Dec}(R)\|_2^2$$

$$E_\theta(R,S|) = -\mathop{\mathbb{E}}_{X,Y} Y \cdot \log\left(\mathrm{Pred}(R)\right) + (1-Y) \cdot \log\left(1-\mathrm{Pred}(R)\right)$$

Censoring Representations with an Adversary: Edwards & Storkey, 2015

RESULTS

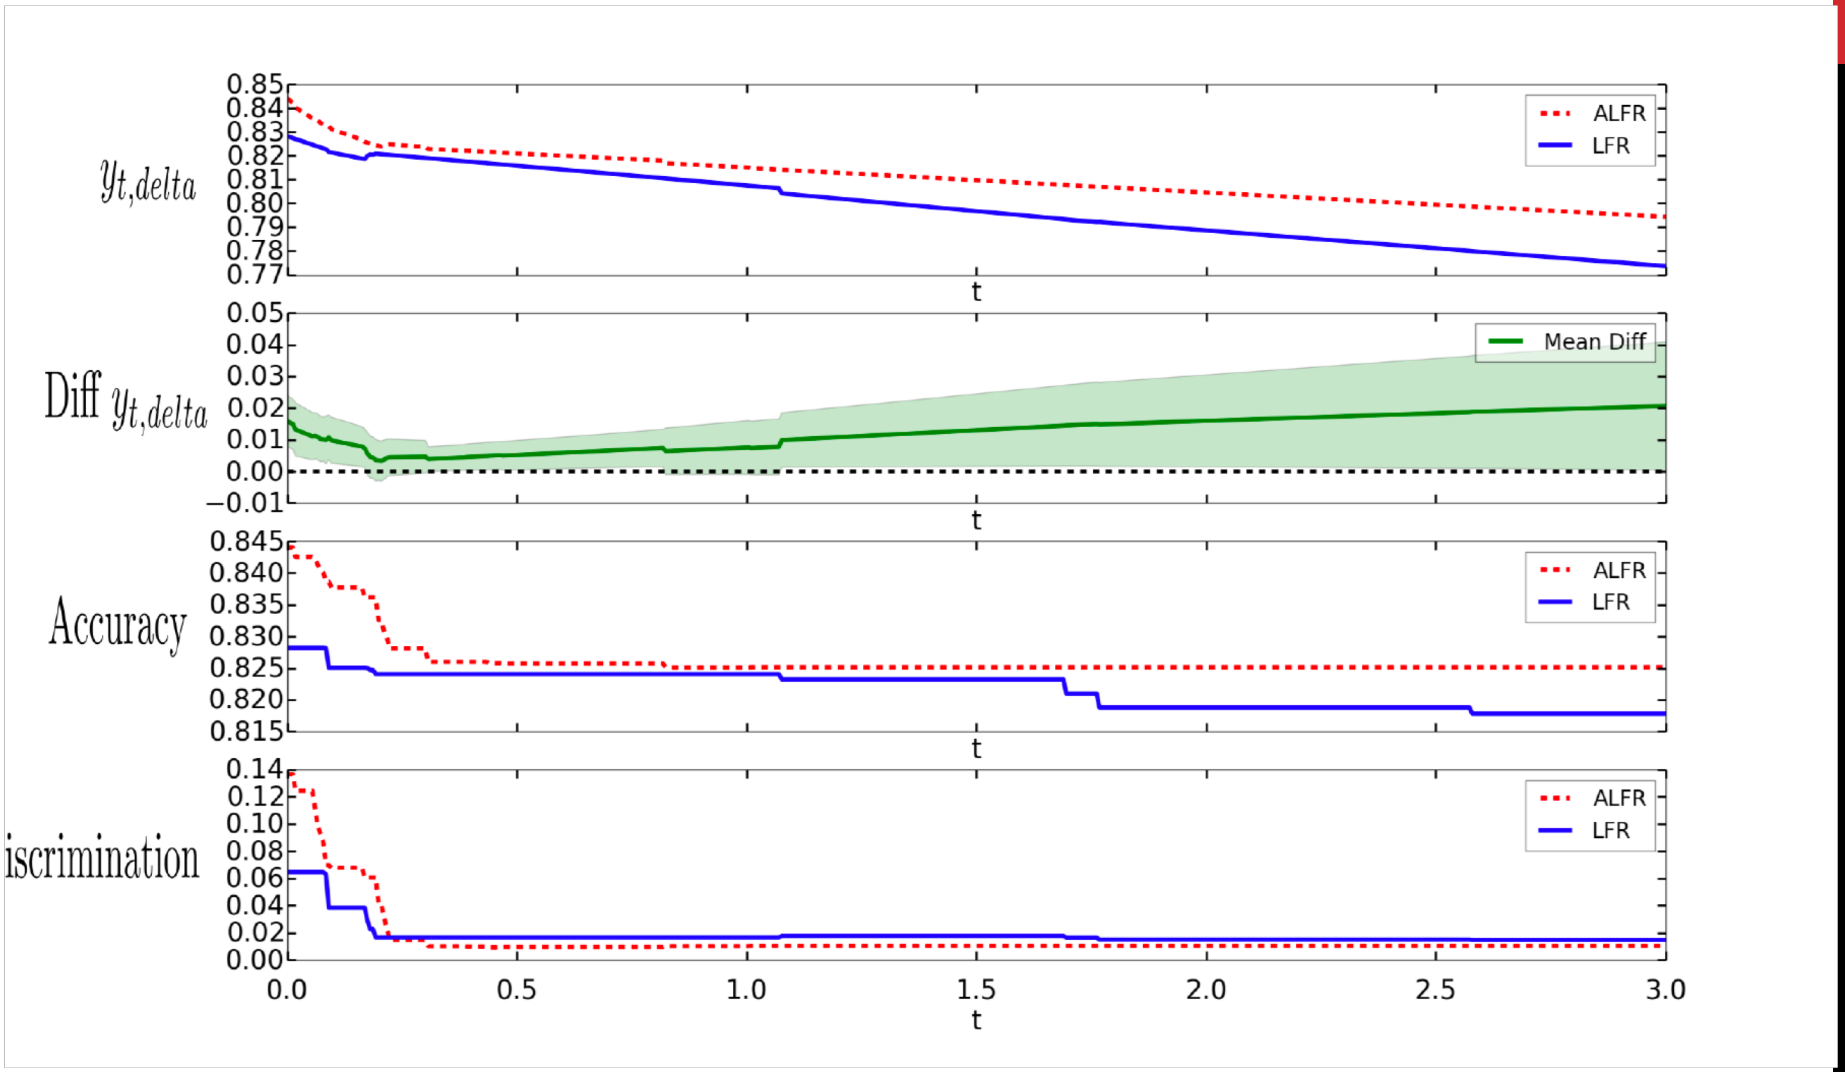
$y_{t,delta}$
0.85
0.84
0.83
0.82
0.81
0.80
0.79
0.78
0.77
t
ALFR
LFR
Diff $y_{t,delta}$
0.05
0.04
0.03
0.02
0.01
0.00
−0.01
t
Mean Diff
Accuracy
0.845
0.840
0.835
0.830
0.825
0.820
0.815
t
ALFR
LFR
iscrimination
0.14
0.12
0.10
0.08
0.06
0.04
0.02
0.00
0.0
0.5
1.0
1.5
2.0
2.5
3.0
t
ALFR
LFR

RESULTS


Davies
Chambers
Doney
Langhans
Scowden
Robertson
Yeakley
Hall
Sherer
Flachs
Smith
May
Spruill
Goulet
Parish
Dennis
Whitfield
Williams
Walton

EQUALIZED ODDS / OPPORTUNITY

Both VFAE and AFLR define fairness as statistical parity

Problems with demographic/statistical parity:

- Coarse measure, not about individuals
- May entail large loss in accuracy

Alternative definition: equal opportunity [Hardt, Price, Srebro, 2016]

- Encourage perfect prediction
- But ensure that the prediction errors are balanced between the groups

$$\Pr\left\{\widehat{Y}=1 \mid A=0, Y=y\right\}=\Pr\left\{\widehat{Y}=1 \mid A=1, Y=y\right\}, \quad y \in \{0,1\}$$

FAIR CLASSIFICATION: DEFINITIONS

Definitions based on predicted outcomes:

- Demographic / statistical parity
- Conditional statistical parity (loan conditioned on credit history, amount, employment)

Definitions based on predicted and actual outcomes:

- Balanced PPV (FDR) – predictive equality
- Balanced FNR (TPR) – equal opportunity
- Balanced FNR and FPR – equalized odds

	Actual – Positive	Actual – Negative
Predicted – Positive	**True Positive (TP)** $\text{PPV} = \frac{TP}{TP+FP}$ $\text{TPR} = \frac{TP}{TP+FN}$	**False Positive (FP)** $\text{FDR} = \frac{FP}{TP+FP}$ $\text{FPR} = \frac{FP}{FP+TN}$
Predicted – Negative	**False Negative (FN)** $\text{FOR} = \frac{FN}{TN+FN}$ $\text{FNR} = \frac{FN}{TP+FN}$	**True Negative (TN)** $\text{NPV} = \frac{TN}{TN+FN}$ $\text{TNR} = \frac{TN}{TN+FP}$

FAIR CLASSIFICATION: MORE DEFINITIONS

Definitions based on predicted probability and actual outcomes:

- Balanced for positive class
- Balance for negative class
- Calibration

Definitions based on similarity

- Individual fairness

FAIR CLASSIFICATION: DEFINITIONS

Most common way to define fair classification is to require some invariance with respect to the sensitive attribute

- Demographic parity: $\hat{Y} \perp A$
- Equalized Odds: $\hat{Y} \perp A|Y$
- Equal Opportunity: $\hat{Y} \perp A|Y = y$, for some y
- Equal (Weak) Calibration: $Y \perp A|\hat{Y}$
- Equal (Strong) Calibration: $Y \perp A|\hat{Y}$ and $\hat{Y} = P(Y = 1)$
- Fair Subgroup Accuracy: $\mathbb{1}[Y = \hat{Y}] \perp A$

Note: Many of these definitions are incompatible!

WHY FAIR REPRESENTATIONS

- Minimize unfair targeting of disadvantaged groups by vendors (worse lines of credit, lower paying jobs)
- Aim: form a data representation that ensures fair classifications downstream

- Consider two types of unfair vendors:
  1. The **indifferent** vendor: does not care about fairness, only maximizes utility
  2. The **malicious** vendor: doesn't care about utility, discriminates unfairly

- Suggests an adversarial learning scheme

LEARNING ADVERSARIALLY FAIR TRANSFERABLE REPRESENTATIONS

Madras, Creager, Pitassi, Zemel, 2018

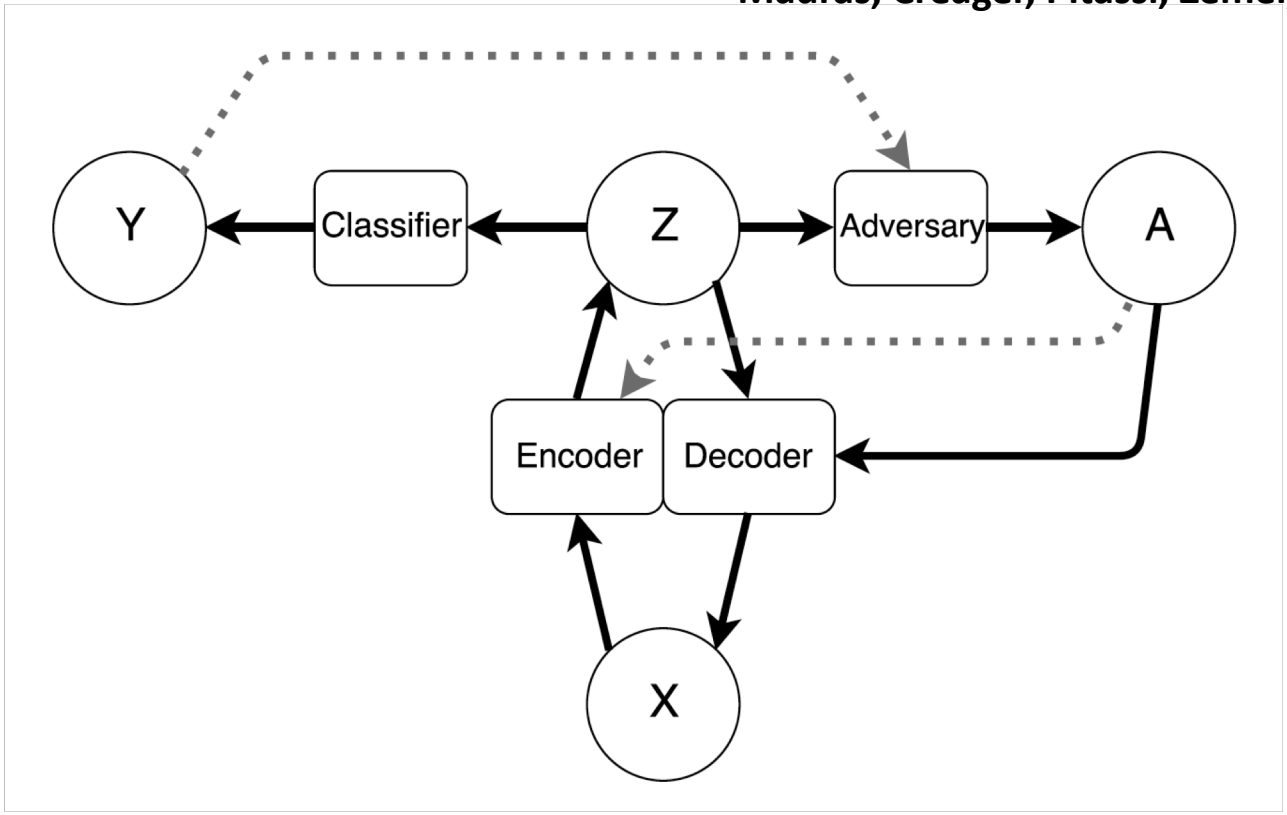

- The classifier is indifferent vendor, forcing the encoder to make the representations useful
- The adversary is the malicious vendor, forcing the encoder to hide the sensitive attributes in the representations

ADVERSARIAL LEARNING IN LAFTR

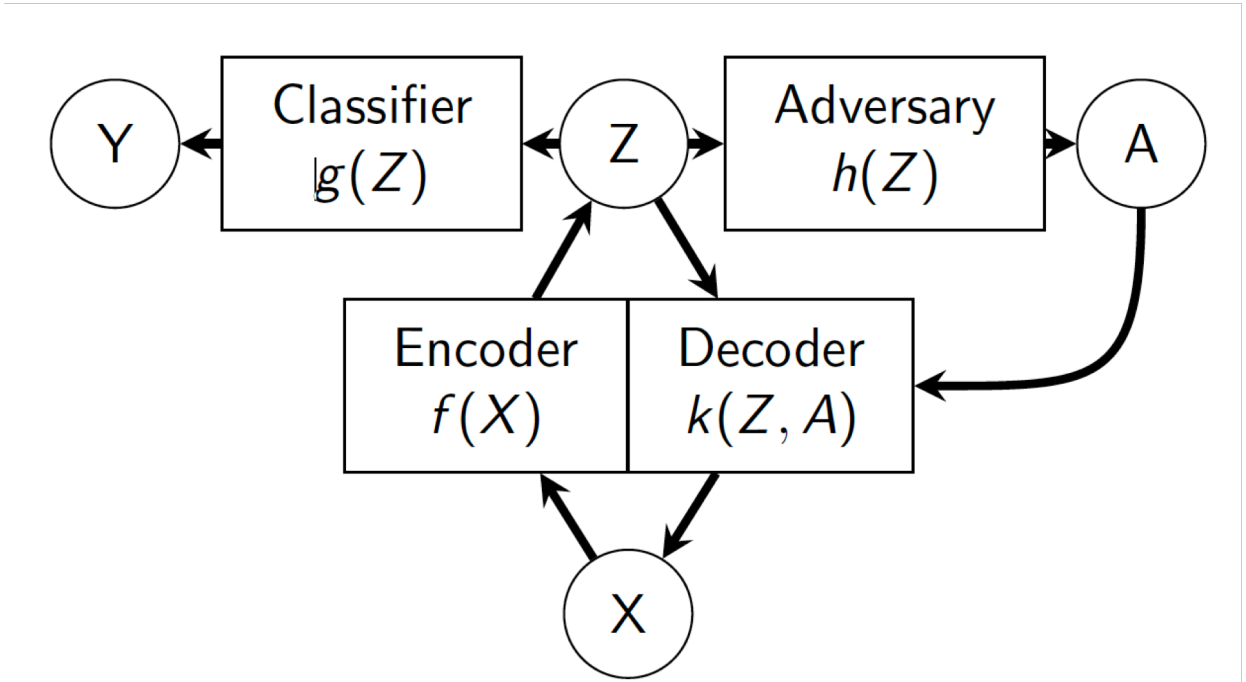

- Our game: encoder-decoder-classifier vs. adversary
- Aim: Learn fair encoder

$$\underset{f,g,k}{\text{minimize}}\ \underset{h}{\text{maximize}}\ \mathbb{E}_{X,Y,A}\left[\mathcal{L}(f,g,h,k)\right]$$

$$\mathcal{L}(f,g,h,k) = \alpha\mathcal{L}_{Class} + \beta\mathcal{L}_{Dec} - \gamma\mathcal{L}_{Adv}$$

ADVERSARIAL OBJECTIVES

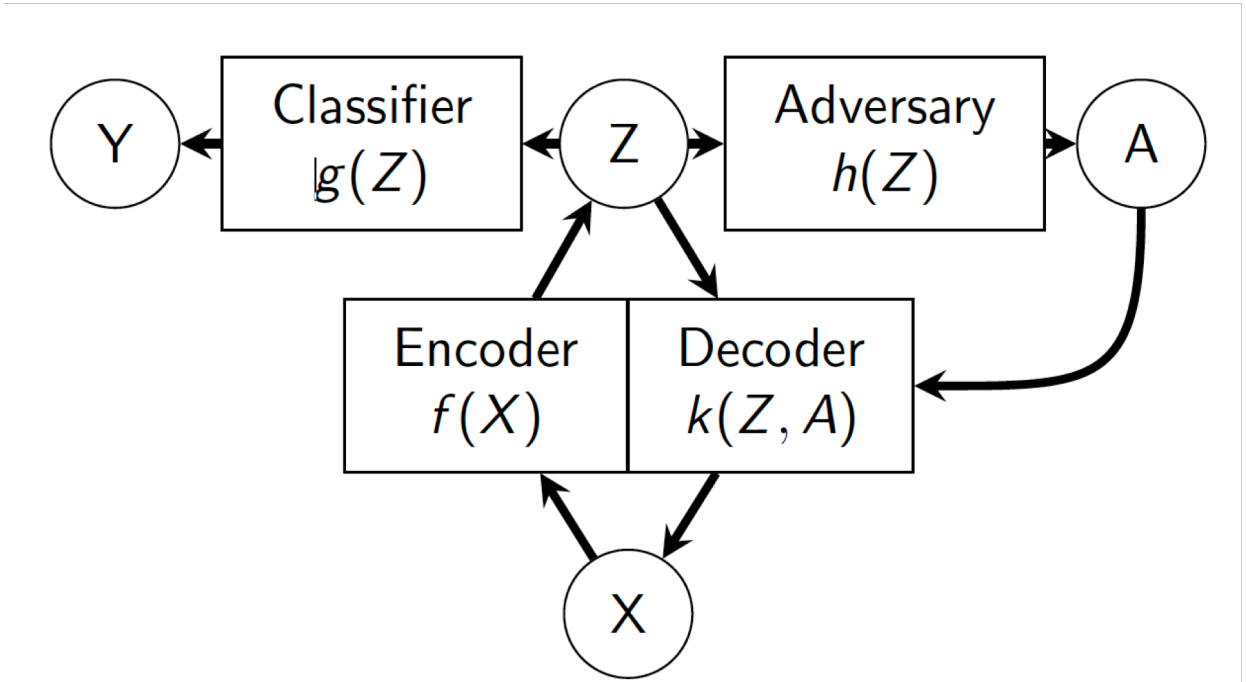

Choice of adversarial objective depends on fairness desideratum

- Demographic parity: $\mathcal{L}_{DP}(h) = \sum_{i \in \{0,1\}} \frac{1}{|\mathcal{D}_i|} \sum_{(x,a) \in \mathcal{D}_i} |h(f(x)) - a|$
- Equalized odds: $\mathcal{L}_{EO}(h) = \sum_{i,j \in \{0,1\}^2} \frac{1}{|\mathcal{D}_i^j|} \sum_{(x,a,y) \in \mathcal{D}_i^j} |h(f(x), y) - a|$
- Equal Opportunity: $\mathcal{L}_{EOpp}(h) = \sum_{i \in \{0,1\}} \frac{1}{|\mathcal{D}_i^1|} \sum_{(x,a) \in \mathcal{D}_i^1} |h(f(x)) - a|$

FROM ADVERSARIAL OBJECTIVES TO FAIRNESS DEFINITIONS

In general: pick the right adversarial loss, encourage the right conditional independencies

- Demographic parity encourages $Z \perp A$ to fool adversary
- Equalized odds encourages $Z \perp A \mid Y$ to fool adversary
- Equal opportunity encourages $Z \perp A \mid Y = 1$ to fool adversary

Note that independencies of $Z = f(x)$ also hold for predictions $\hat{Y} = g(Z)$

We show: In the adversarial limit, these objectives guarantee these fairness metrics!

- The key is to connect predictability of A by the adversary $h(Z)$ to unfairness in the classifier $g(Z)$

EXPERIMENTS

Datasets

1. Adult Income

Size: 45,222 instances, 14 attributes

Task: predict whether or not annual income > 50K

Sensitive feature: Gender

2. Heritage Health

Size: 147,473 instances, 139 attributes

Task: predict patient's Charlson Index (co-morbidity)

Sensitive feature: Age

Models

Encoder, classifier, adversary: each single hidden-layer MLP (8; 20 hidden units)

RESULTS: FAIR CLASSIFICATION

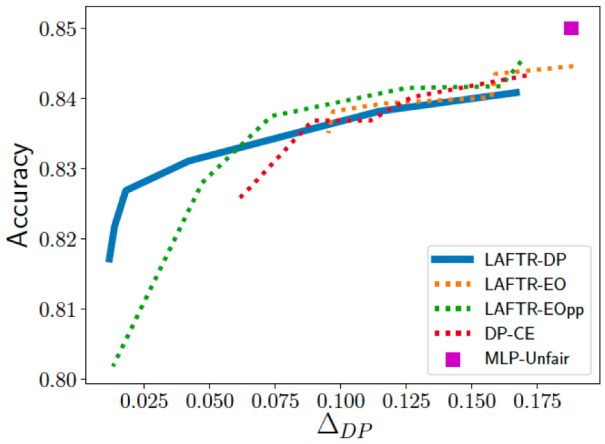

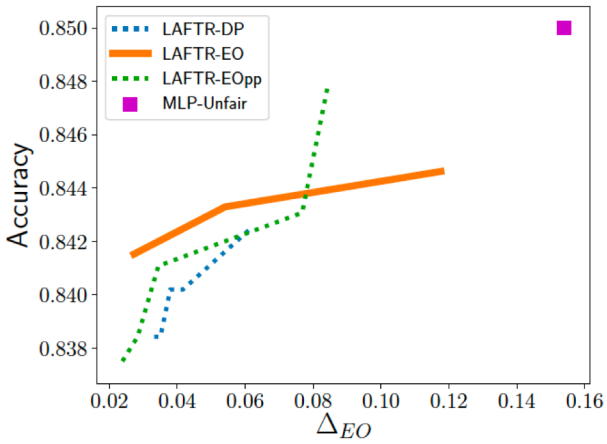

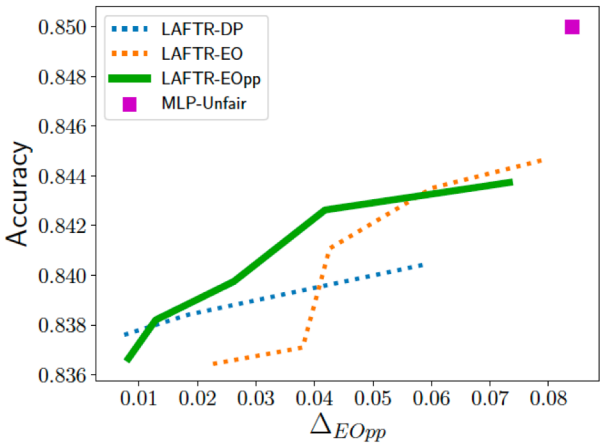

- Train with 2-step process to simulate owner → vendor framework
- Tradeoffs between accuracy and fairness metrics produced by different LAFTR loss functions
- Achieves best solutions, wrt fairness-accuracy tradeoff

RESULTS: FAIRNESS METRICS

METHOD	Δ_{DP}	Δ_{EO}	Δ_{EOpp}	ACC.
MLP (UNFAIR)	0.381	0.476	0.231	**0.785**
LAFTR-EO	0.152	**0.050**	0.036	0.763
	0.143	**0.052**	0.032	0.752
LAFTR-EOPP	0.087	0.092	**0.010**	0.742
	0.113	0.063	**0.024**	0.735
LAFTR-DP	**0.041**	0.140	0.025	0.731
	0.002	0.196	0.031	0.728

SETUP: FAIR TRANSFER LEARNING

- Downstream vendors will have unknown prediction tasks
- Does fairness transfer?
- We test this as follows:
  1. Train encoder f on data X, with label Y
  2. Freeze encoder f
  3. On new data X', train classifier on top of $f(X')$, with new task label Y'
  4. Observe fairness and accuracy of this new classifier on new task Y'
- Compare LAFTR encoder f to other encoders
- We use Heritage Health dataset
  - Y is Charlson comorbidity index > 0
  - Y' is whether or not a certain type of insurance claim was made
  - Check for fairness w.r.t. age

RESULTS : FAIR TRANSFER LEARNING

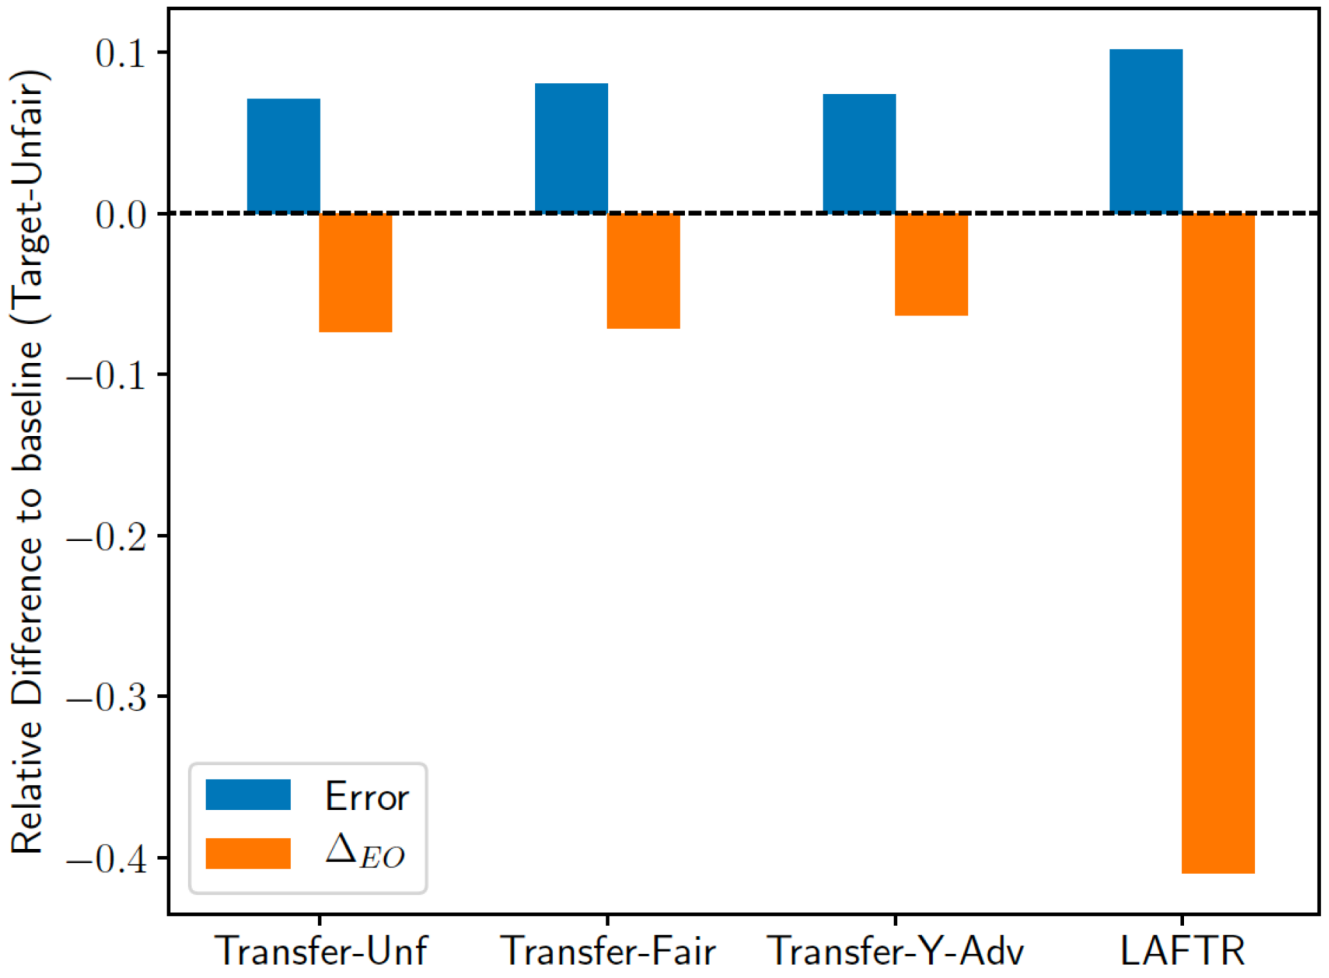

Fair transfer learning on Health dataset. Down is better in both metrics.