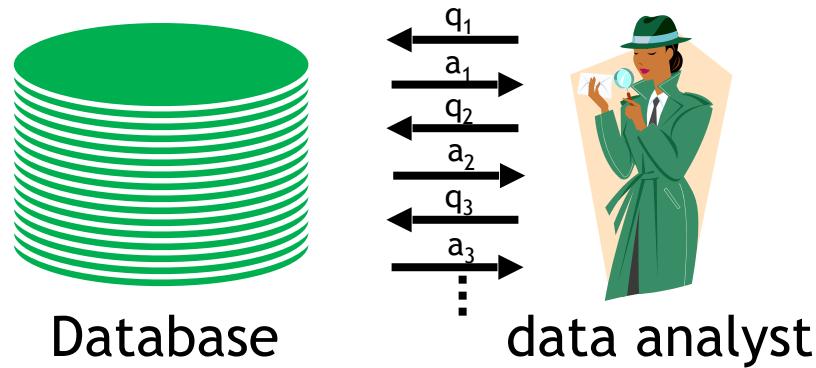


Privacy-Preserving Data Analysis



- ▶ Census, epidemic detection based on OTC drug purchases; analysis of loan application data for evidence of discrimination,....
- ▶ 50+ year old problem

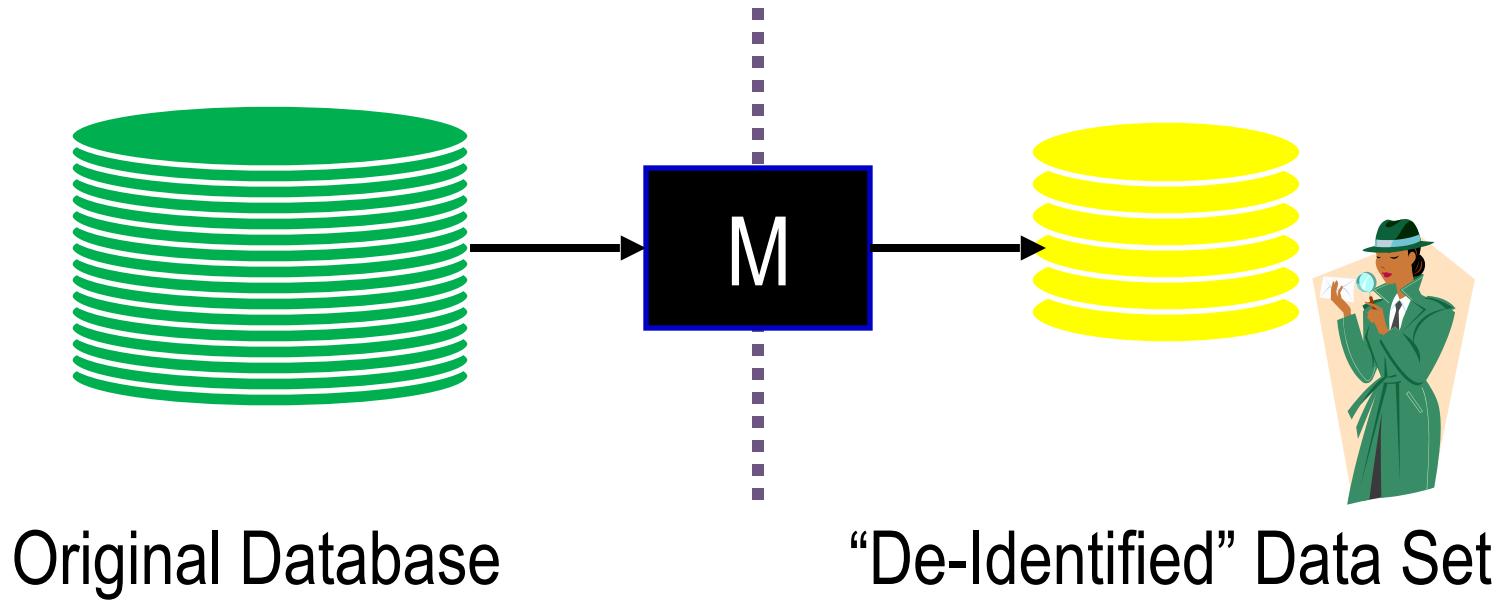
What analyses on a database might violate privacy? What analyses are privacy-preserving?

what to promise?

delete identifying information

maybe not

“De-Identification”?



De-Identified data isn’t.



Latanya Sweeney's Attack (1997)

Massachusetts hospital discharge dataset

• **Medical Data Released as Anonymous**

SSN	Name	City	Date Of Birth	Sex	ZIP	Marital Status	Problem
			09/27/64	female	02139	divorced	hypertension
			09/30/64	female	02139	divorced	obesity
	asian		04/18/64	male	02139	married	chest pain
	asian		04/15/64	male	02139	married	obesity
	black		03/13/63	male	02138	married	hypertension
	black		03/18/63	male	02138	married	shortness of breath
	black		09/13/64	female	02141	married	shortness of breath
	black		09/07/64	female	02141	married	obesity
	white		05/14/61	male	02138	single	chest pain
	white		05/08/61	male	02138	single	obesity
	white		09/15/61	female	02142	widow	shortness of breath

• **Voter List**

Name	Address	City	ZIP	DOB	Sex	Party
.....
.....
Sue J. Carlson	1459 Main St.	Cambridge	02142	9/15/61	female	democrat

Figure 1: Re-identifying anonymous data by linking to external data

Public voter dataset

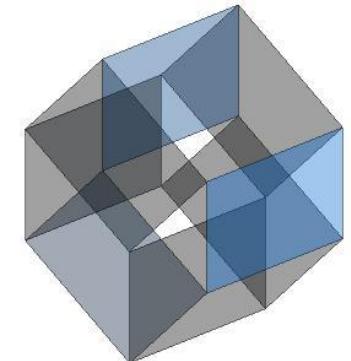
K-Anonymity: Intuition

- The information for each person contained in the released table cannot be distinguished from at least $k-1$ individuals whose information also appears in the release
 - Example: you try to identify a man in the released table, but the only information you have is his birth date and gender. There are k men in the table with the same birth date and gender.
- Any quasi-identifier present in the released table must appear in at least k records

Curse of Dimensionality

Aggarwal (VLDB 2005)

- Generalization fundamentally relies on **spatial locality**
 - Each record must have k close neighbors
- Real-world datasets are **very sparse**
 - Many attributes (dimensions)
 - Netflix Prize dataset: 17,000 dimensions
 - Amazon customer records: several million dimensions
 - “Nearest neighbor” is **very far**
- Projection to low dimensions loses all info \Rightarrow **k -anonymized datasets are useless**



what to promise?

only ask questions that pertain
to large populations

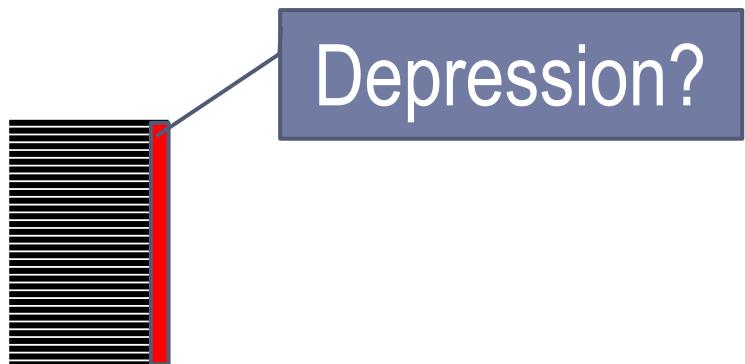
maybe not

The Statistics Masquerade

- ▶ Differencing Attack
 - ▶ *How many members of House of Representatives have sickle cell trait?*
 - ▶ *How many members of House, other than the Speaker, have the trait?*
- ▶ Needle in a Haystack
 - ▶ Determine presence of an individual's genomic data in GWAS case group



- ▶ The Big Bang attack
 - ▶ Reconstruct “depression” bit column



Fundamental Law of Info Recovery

- ▶ “Overly accurate” estimates of “too many” statistics is blatantly non-private.



what to promise?

access to the output should not enable one to learn anything about an individual that could not be learned without access

is this possible?

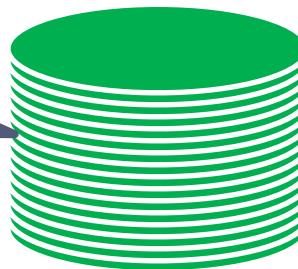
hint: either privacy or utility separately is easy

what to promise?

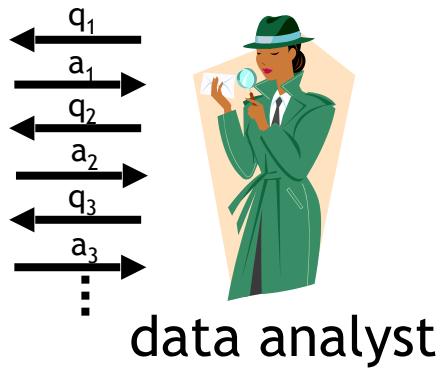
access to the output should not enable one to learn anything about an individual that could not be learned without access

is this
desirable?

Privacy-Preserving Data Analysis?



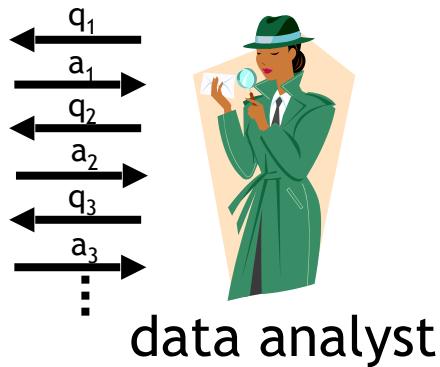
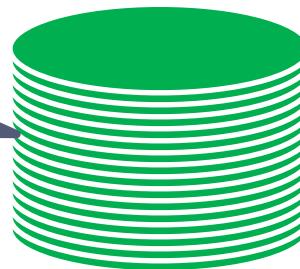
Database



- ▶ “Can’t learn anything new about Helen”?
- ▶ Dalenius, 1977; Goldwasser and Micali: semantic security 1982



Privacy-Preserving Data Analysis?



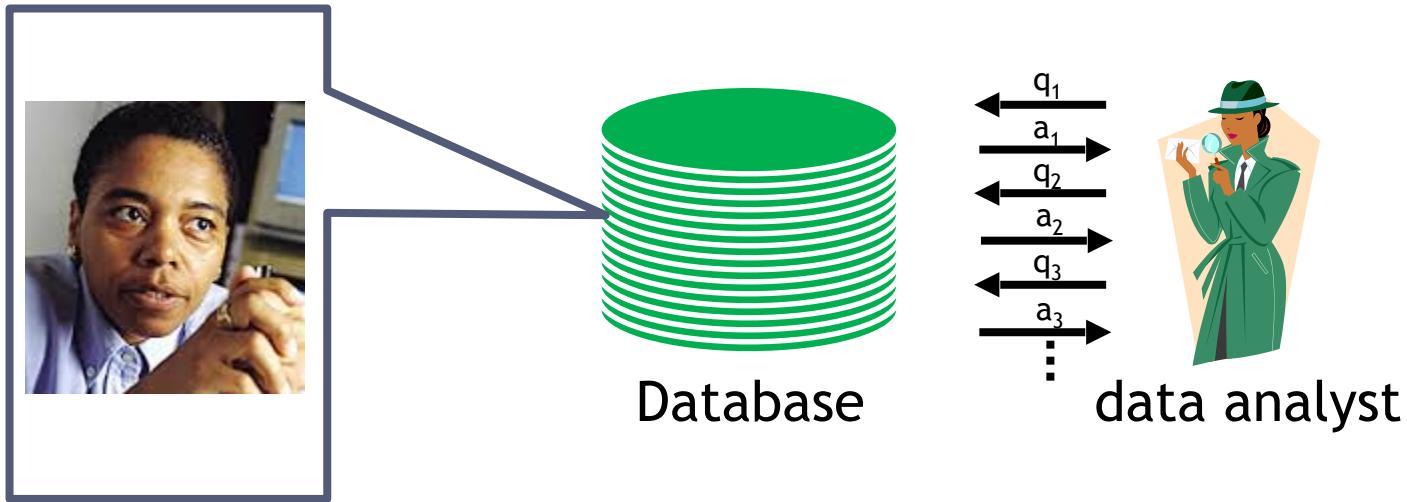
- ▶ “Can’t learn anything new about Helen”?
- ▶ Then what is the point?



what to promise?

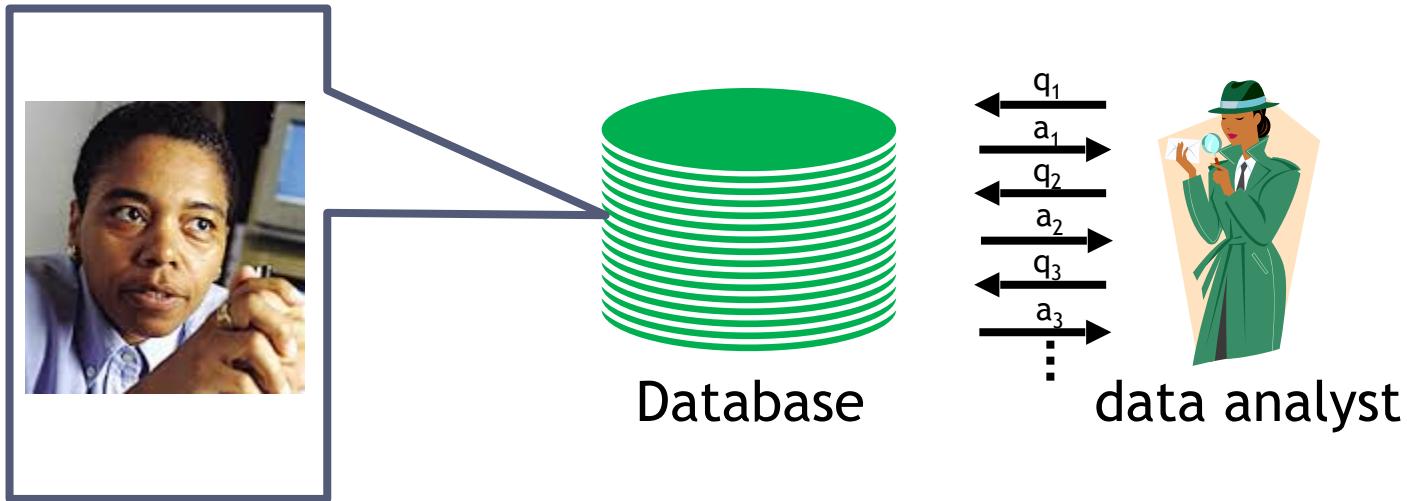
access to the output should not enable one to learn much more about an individual than could be learned via the same analysis omitting that individual from the database

Privacy-Preserving Data Analysis?



- ▶ Ideally: learn same things if Helen is replaced by another random member of the population

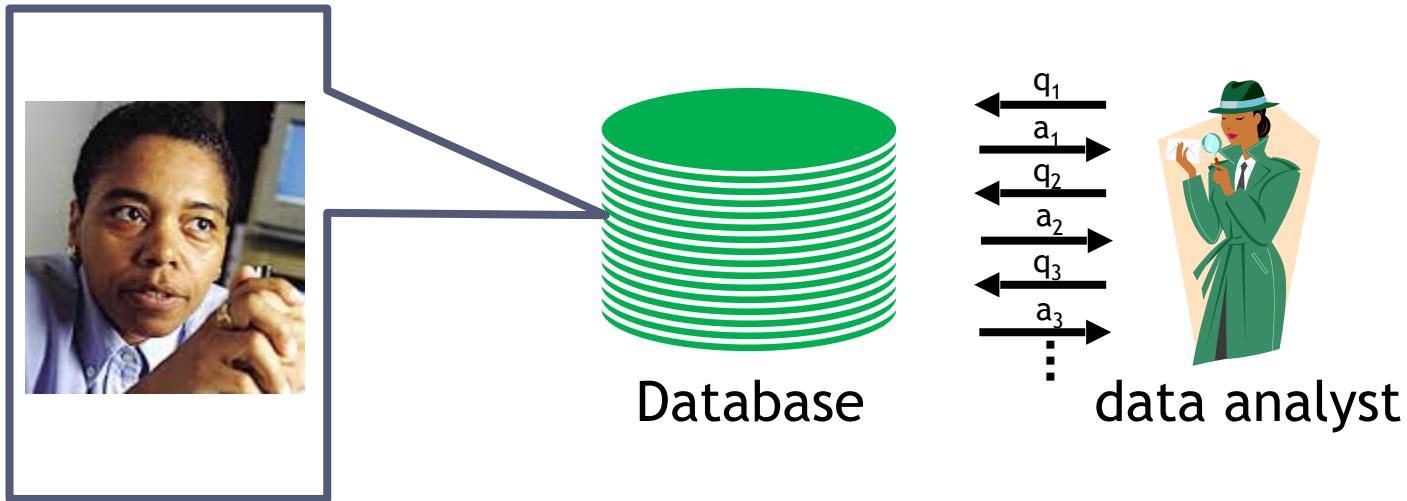
Privacy-Preserving Data Analysis?



- ▶ Ideally: learn same things if Helen is replaced by another random member of the population (“stability”)



Privacy-Preserving Data Analysis?



- ▶ Stability preserves Helen's privacy AND prevents over-fitting
- ▶ Privacy and Generalization are aligned!



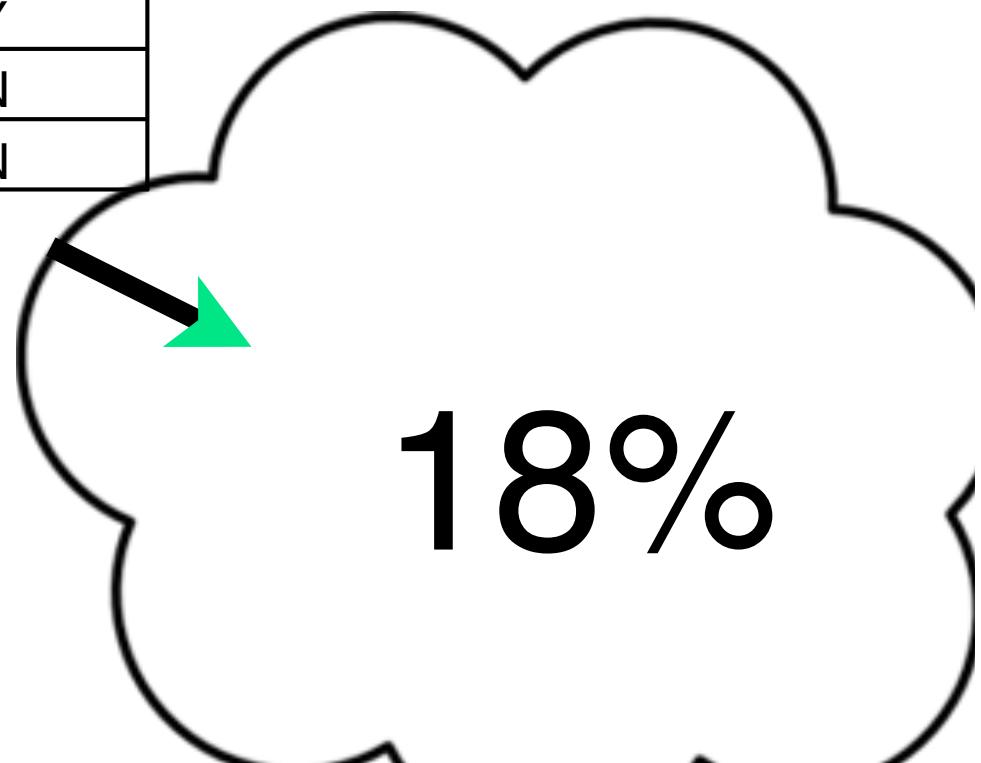
Differential Privacy

- ▶ The outcome of any analysis is essentially equally likely, independent of whether any individual joins, or refrains from joining, the dataset.
 - ▶ Helen goes away, Latanya joins, Helen is replaced by Latanya

what to promise?

think of output as randomized

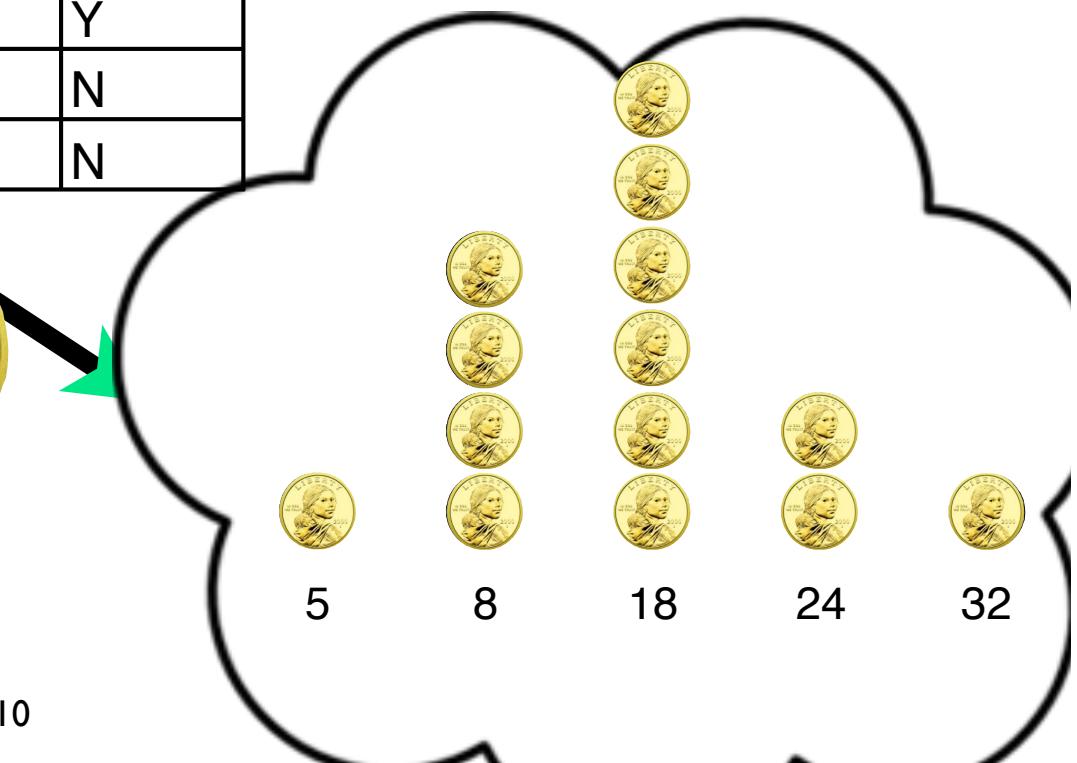
name	DOB	sex	weight	smoker	lung cancer
John Doe	12/1/51	M	185	Y	N
Jane Smith	3/3/46	F	140	N	N
Ellen Jones	4/24/59	F	160	Y	Y
Jennifer Kim	3/1/70	F	135	N	N
Rachel Waters	9/5/43	F	140	N	N



what to promise?

think of output as randomized

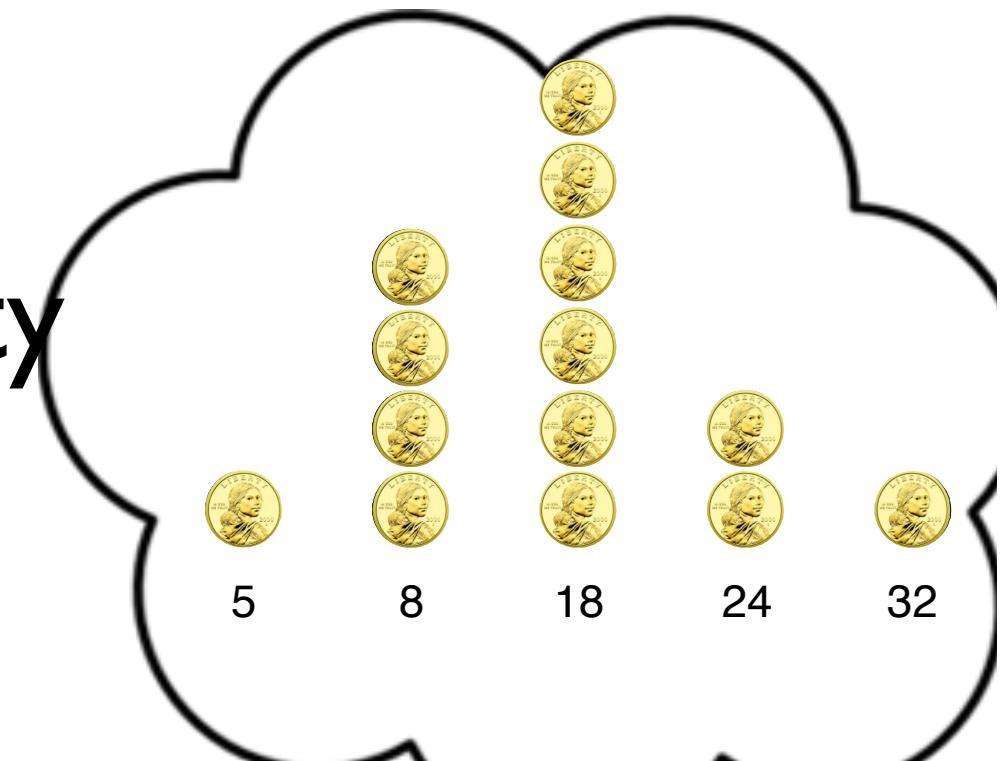
name	DOB	sex	weight	smoker	lung cancer
John Doe	12/1/51	M	185	Y	N
Jane Smith	3/3/46	F	140	N	N
Ellen Jones	4/24/59	F	160	Y	Y
Jennifer Kim	3/1/70	F	135	N	N
Rachel Waters	9/5/43	F	140	N	N



what to promise?

think of output as randomized

promise: if you leave
the database, no
outcome will
change probability
by very much



statistical database model

X set of possible entries/rows

one row per person

database x a set of rows; $x \in \mathbb{N}^{|X|}$
(histogram)

name	DOB	sex	weight	smoker	lung cancer
John Doe	12/1/51	M	185	Y	N
Jane Smith	3/3/46	F	140	N	N
Ellen Jones	4/24/59	F	160	Y	Y
Jennifer Kim	3/1/70	F	135	N	N
Rachel Waters	9/5/43	F	140	N	N

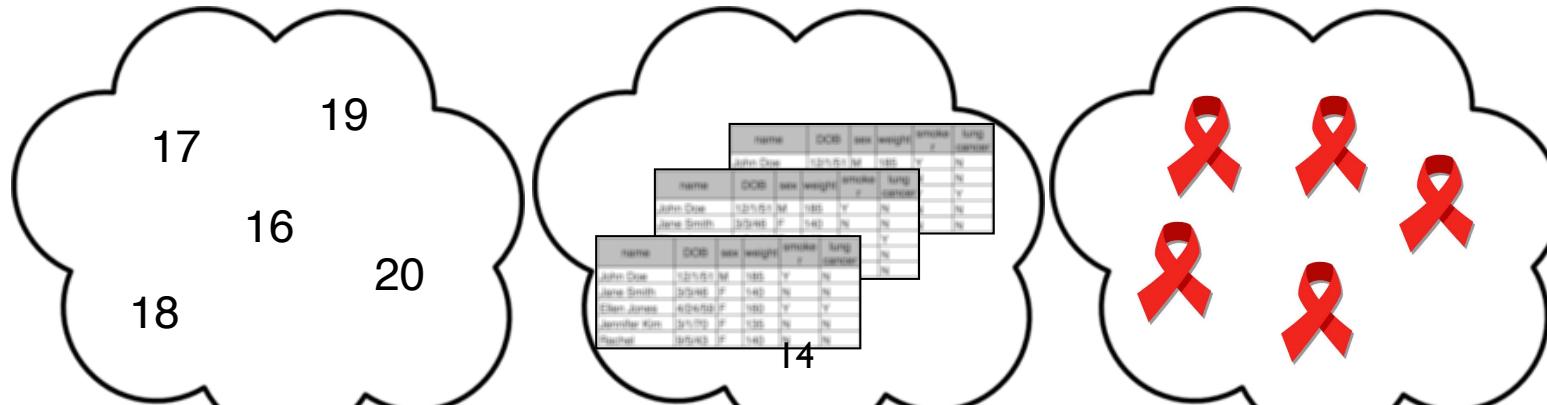
analyst objective

wishes to compute on $x \in \mathbb{N}^{|X|}$

fit a model, compute a statistic, share
“sanitized” data

preserve privacy of individuals

design randomized algorithm M mapping x to into
outcome space, that masks small changes in x



neighboring databases

what's a small change?

require nearly identical behavior on neighboring databases differing by the addition or removal of a single row:

$$\|x - y\|_1 \leq 1$$

for $x, y \in \mathbb{N}^{|X|}$

differential privacy

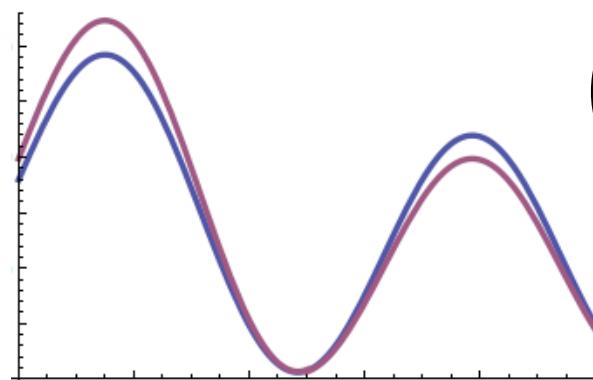
[DinurNissim03, DworkNissimMcSherrySmith06, Dwork06]

ϵ -Differential Privacy for algorithm M :

for any two neighboring data sets x_1, x_2 , differing by the addition or removal of a single row

any $S \subseteq \text{range}(M)$,

$$\Pr[M(x_1) \in S] \leq e^\epsilon \Pr[M(x_2) \in S]$$

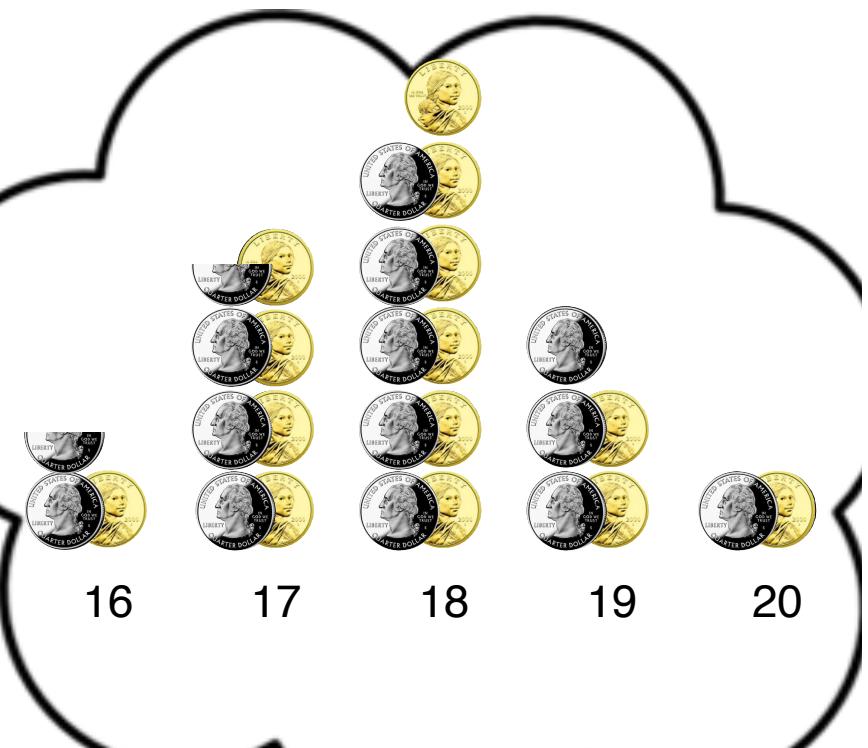


$$e^\epsilon \sim (1 + \epsilon)$$

differential privacy

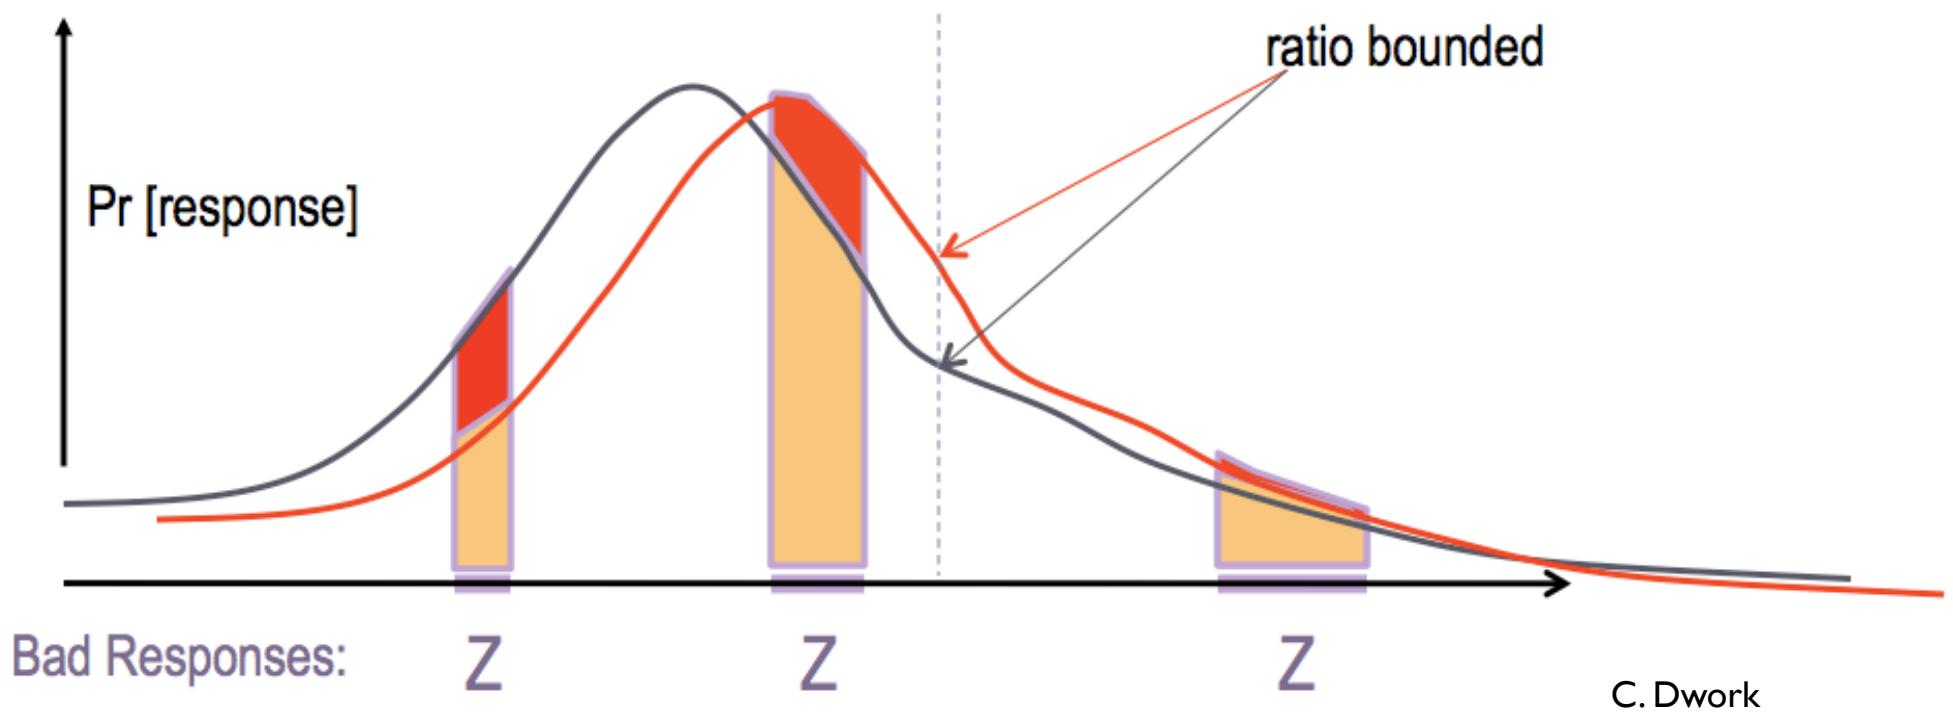
$$\Pr[M(x_1) \in S] \leq e^\varepsilon \Pr[M(x_2) \in S]$$

name	DOB	sex	weight	smoker	lung cancer
John Doe	12/1/51	M	185	Y	N
Jane Smith	3/3/46	F	140	N	N
Ellen Jones	4/24/59	F	160	Y	Y
Jennifer Kim	3/1/70	F	135	N	N
Rachel Waters	9/5/43	F	140	N	N



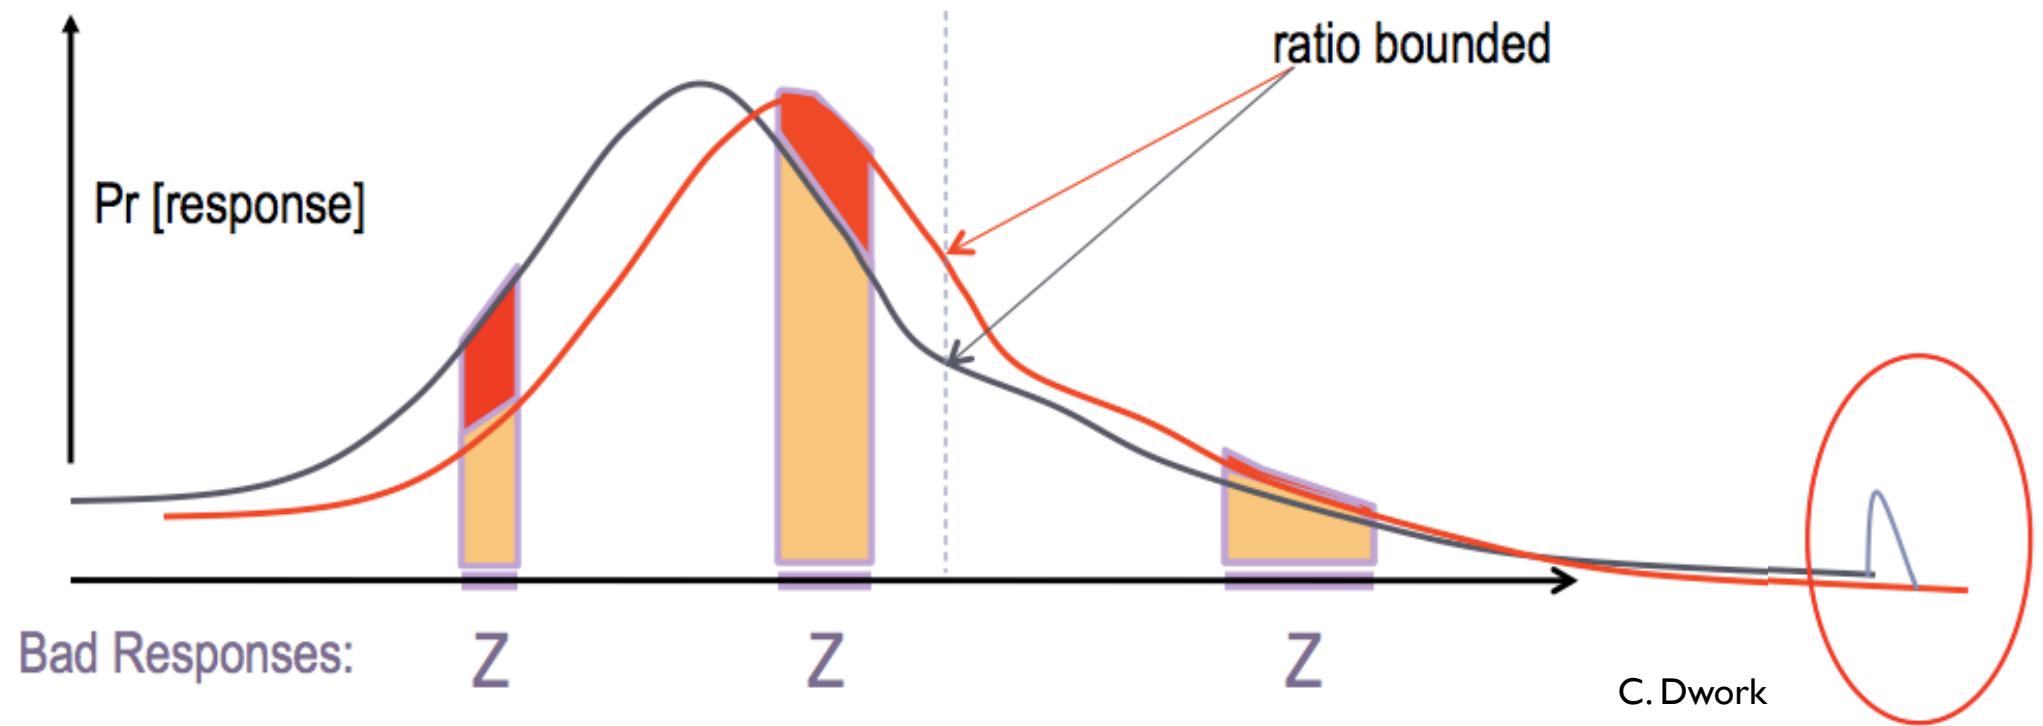
differential privacy

$$\Pr[M(x_1) \in S] \leq e^\varepsilon \Pr[M(x_2) \in S]$$



(ε, δ) -differential privacy

$$\Pr[M(x_1) \in S] \leq e^\varepsilon \Pr[M(x_2) \in S] + \delta$$



differential privacy

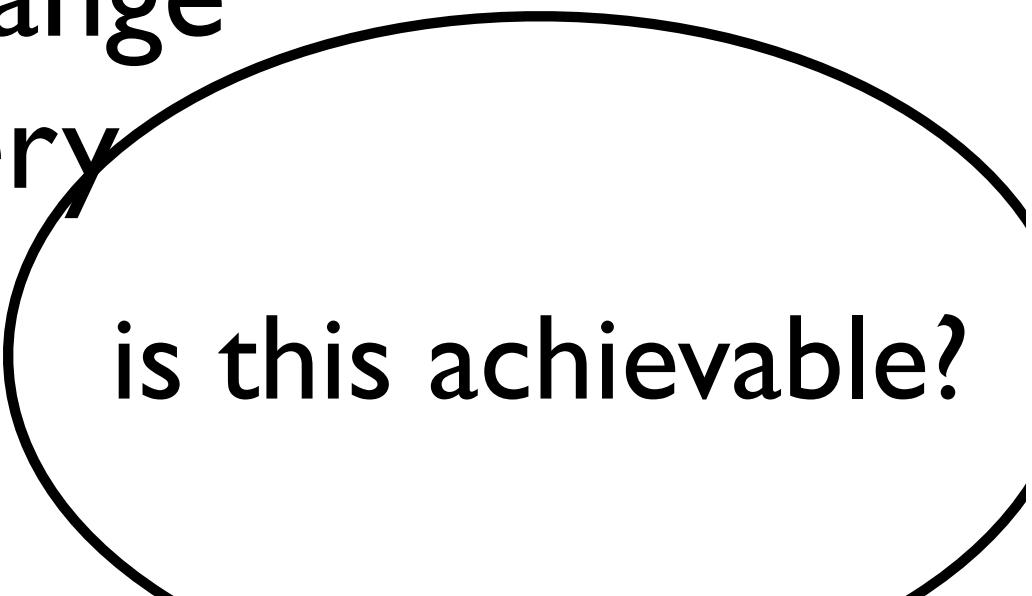
$$\Pr[M(x_1) \in S] \leq e^\varepsilon \Pr[M(x_2) \in S]$$

Is a statistical property of mechanism behavior
unaffected by auxiliary information
independent of adversary's computational
power

differential privacy

$$\Pr[M(x_1) \in S] \leq e^\varepsilon \Pr[M(x_2) \in S]$$

promise: if you leave
the database, no
outcome will change
probability by very
much



is this achievable?

yes!





Differential privacy



Apple will not see your data



today

Formalizing privacy

→ Privacy properties and basic tools

Randomized Response

differential privacy

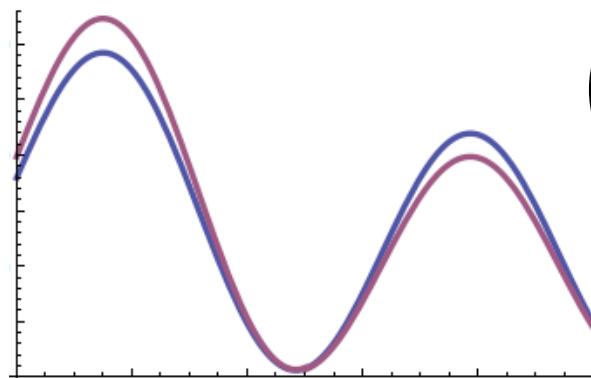
[DinurNissim03, DworkNissimMcSherrySmith06, Dwork06]

ϵ -Differential Privacy for algorithm M :

for any two neighboring data sets x_1, x_2 , differing by the addition or removal of a single row

any $S \subseteq \text{range}(M)$,

$$\Pr[M(x_1) \in S] \leq e^\epsilon \Pr[M(x_2) \in S]$$



$$e^\epsilon \sim (1 + \epsilon)$$

group privacy

Thm. Any $(\varepsilon, 0)$ -DP mechanism M is $(k \varepsilon, 0)$ -DP for groups of size k . i.e., for all

$$||x - y||_1 \leq k$$

and any $S \subseteq \text{range}(M)$,

$$\Pr[M(x) \in S] \leq e^{\varepsilon k} \Pr[M(y) \in S]$$

post-processing

Thm. Let $M : \mathbb{N}^{|X|} \rightarrow R$ be (ε, δ) -DP.

Let $f : R \rightarrow R'$ be an arbitrary randomized mapping.

Then $f \circ M : \mathbb{N}^{|X|} \rightarrow R'$ is (ε, δ) -DP.

composition

[DworkKenthapadiMcSherryMironovNaor06,DworkLei09]

Thm. For $i \in [k]$, let $M_i : \mathbb{N}^{|X|} \rightarrow R_i$ be $(\varepsilon_i, \delta_i)$ -DP. Then the mechanism $(M_1(x), \dots, M_k(x))$ is $(\sum_i \varepsilon_i, \sum_i \delta_i)$ -DP.

actually, holds even if subsequent computations chosen as function of previous results

“advanced” version

Is bigger delta better for privacy, or worse?
What about epsilon?

today

Formalizing privacy

Privacy properties and basic tools

→ Randomized Response

Randomized Response

[Warner65]

flip a coin

if tails, respond truthfully

if heads, flip a second coin and respond
“yes” if heads; respond “no” if tails

Claim. Randomized Response is $(\ln 3, 0)$ -DP.

Proof.
$$\frac{\Pr[\text{Response} = \text{Yes} | \text{Truth} = \text{Yes}]}{\Pr[\text{Response} = \text{Yes} | \text{Truth} = \text{No}]}$$

$$= \frac{3/4}{1/4} = \frac{\Pr[\text{Response} = \text{No} | \text{Truth} = \text{No}]}{\Pr[\text{Response} = \text{No} | \text{Truth} = \text{Yes}]} = 3.$$

Randomized Response

Given database $x = x_1, \dots, x_n$ where $x_i \in \{0, 1\}$

(say $x_i = 1$ if person committed crime
 $= 0$ if person did not commit crime)

Query : $\sum_i x_i / n$ (= fraction of people that committed crime)

Mechanism :

Step 1. For $i=1 \dots n$

Let $y_i = x_i$ with probability $\frac{3}{4}$
 $y_i = 1-x_i$ with probability $\frac{1}{4}$

Step 2. Let $f(y_1 \dots y_n) = \sum_i y_i / n$
Output $f(y_1 \dots y_n)$

Lemma Mechanism is $(\ln 3, 0)$ -dp

pt First we show that the output of step 1

$y_1 \dots y_n$ is $(\ln 3, 0)$ -dp. Then by

post processing, $f(y_1 \dots y_n)$ is also $(\ln 3, 0)$ -dp.

Consider 2 neighboring databases

$$\begin{aligned} x &= x_1 x_2 \dots x_n \\ x' &= x_1 \dots \bar{x}_i \dots x_n \end{aligned} \quad \left. \begin{array}{l} \text{differ only on} \\ \text{coord. } i \end{array} \right.$$

Show $\forall y_1 \dots y_n \frac{\Pr(y_1 \dots y_n | x_1 \dots x_n)}{\Pr(y_1 \dots y_n | x_1 \dots \bar{x}_i \dots x_n)} = 3$

$$\frac{\Pr(y_1 \dots y_n | x_1 \dots x_n)}{\Pr(y_1 \dots y_n | x_1 \dots \bar{x}_i \dots x_n)} = \frac{\Pr(y_1 | x_1) \cdot \Pr(y_2 | x_2) \dots \Pr(y_n | x_n)}{\Pr(y_1 | x_1) \cdot \Pr(y_i | \bar{x}_i) \cdot \Pr(y_n | x_n)} = \frac{\Pr(y_i | x_i)}{\Pr(y_i | \bar{x}_i)}$$

$$\frac{\Pr(Y_i | X_i)}{\Pr(Y_i | \bar{X}_i)} = \frac{\frac{3}{4}}{\frac{1}{4}} = 3 \geq e^{\varepsilon}$$

$$\Rightarrow \varepsilon \leq \ln 3$$

accuracy of Randomized Response:

Let $n' = \#$ of respondents who say 1 ($= \sum_i Y_i$)

$$\text{let } p = \sum_i X_i / n$$

$$E(n') = (pn) \frac{3}{4} + (1-p)n \frac{1}{4} = p\frac{n}{2} + \frac{1}{4}$$

$$\text{So max likelihood estimator of } p, \hat{p} \text{ is } (n' - \frac{1}{4}) \frac{2}{n} = \frac{2n'}{n} - \frac{1}{2}$$

$$\text{and variance of } \hat{p} = \frac{p(1-p)}{n} + \frac{\frac{3}{4}(\frac{1}{4})}{n(2 \cdot \frac{3}{4} - 1)^2} = \frac{p(1-p)}{n} + \frac{3}{4n}$$